Exact certificates for quantum ordered search: seven queries for 700,000 elements and six for 91,000

Claude Opus 5.5 Extra Effort[1]

[1] *Claude Opus 5.5 Extra Effort*

(Dated: October 31, 2026)

An exact quantum algorithm that searches an ordered list of n elements with k queries, applied recursively, searches N elements with $k\lceil\log_n N\rceil$ queries, so each explicit pair (k,n) gives the bound $(k/\log_2 n)\log_2 N$. Carolan and Childs recently showed that the optimal constant is $\ln 2/\pi \approx 0.2206$, with parameters chosen for sufficiently large n and no explicit finite instance. The explicit finite algorithms known so far come from numerical optimization, and to our knowledge none had been checked in exact arithmetic. We give exact integer certificates for nine explicit algorithms, each checked by a standard-library verifier and by an independent checker built on FLINT. Seven queries search 700,000 elements, which gives $0.360509\log_2 N$ queries. This is the best explicit exact ordered-search algorithm known to us: it improves on the numerical six-query algorithm for 90,000 elements of Wu *et al.* (0.364572) and on the five-query algorithm for 7265 elements (0.389810). Six queries search 91,000 elements, so with the infeasibility certificate of Wu *et al.* the largest list that six queries can search exactly satisfies $91000 \le N^*(6) < 94000$. We also certify Wu *et al.*'s six queries for 90,000 elements and four queries for 605 elements (the optimum found by Childs, Landahl and Parrilo) exactly. A certificate is a chain of symmetric Laurent polynomials in the characterization of Farhi, Goldstone, Gutmann and Sipser, with spectral factors that turn nonnegativity into a diagonal-dominance check. The chains were found by a smoothed max–min optimization over the values of the polynomials at roots of unity, where the chain conditions hold by construction. At $n = 700000$ the roots of the shifted polynomials come within 8.7×10^{-9} of the unit circle, and the spectral factors needed a local quadrature correction. The search, the certificates and this manuscript were produced by a language-model agent.

I. INTRODUCTION

In the ordered search problem the input is a sorted list of n elements and a target, and the task is to find the position of the target. Equivalently, the input is a hidden string $x = 0^j1^{n-j}$ and the algorithm must output j, using queries to the bits of x. Binary search uses $\lceil\log_2 n\rceil$ queries and is optimal classically. A quantum computer can do better by a constant factor, and the size of that factor has been studied for more than 25 years.

Lower bounds. After bounds by Buhrman and de Wolf [1], Farhi *et al.* [2] and Ambainis [3], Høyer, Neerbek and Shi proved that every exact quantum algorithm needs more than $(\ln n - 1)/\pi$ queries, about $0.2206\log_2 n$ [4]. Childs and Lee showed that the adversary method cannot give more, up to an additive constant [5].

Upper bounds. Farhi, Goldstone, Gutmann and Sipser (FGGS) introduced translation-invariant algorithms and found an exact three-query algorithm for $n = 52$, which gives $3\lceil\log_{52} N\rceil \approx 0.526\log_2 N$ queries [6]. Brookes, Jacokes and Landahl reached $4\log_{550} N \approx 0.439\log_2 N$ by gradient descent (as cited in Refs. [7, 9]), and Childs, Landahl and Parrilo (CLP) wrote the FGGS conditions as a semidefinite program and found the four-query optimum $n = 605$, $0.433\log_2 N$ [7]. Ben-Or and Hassidim gave a zero-error algorithm with about $0.32\log_2 N$ expected queries [8]. Carolan, Childs, Kovacs-Deak and Schaeffer (CCKS) proved that translation invariance loses nothing, characterized exact k-query algorithms by a polynomial program, and found the five-query optimum $n = 7265$ $(0.390\log_2 N)$ with a linear-programming relaxation [9]. In 2026, Wu *et al.* reported "strong numerical evidence" from a matrix-free GPU semidefinite program that six queries search 90000 elements $(0.3646\log_2 N)$, and certified that they cannot search 94000 [10]. Shortly afterwards Carolan and Childs settled the asymptotic question: exact algorithms with $(1/\pi)\ln n + o(\log n)$ queries exist [11]. Their construction is analytic, its parameters are chosen for sufficiently large n, and it gives no explicit finite instance.

So the asymptotic constant is known, and what remains open is finite: the largest $n = N^*(k)$ that k queries can search exactly. Every certified pair (k,n) gives an explicit algorithm with $(k/\log_2 n)\log_2 N + k$ queries. The largest values known before this work were $n = 6$, 56, 605 and 7265 for $k = 2,\dots,5$, and $90000 \lesssim N^*(6) < 94000$, all found numerically. To our knowledge none of the algorithms behind them had been checked in exact arithmetic.

This paper gives exact integer certificates for nine explicit algorithms (Tables I and III). The main results are:

- **Seven queries search 700,000 elements**, so ordered search on N elements needs at most $7\lceil\log_{700000} N\rceil \approx 0.360509\log_2 N$ queries. Its constant is 1.1% below that of the numerical $(6, 90000)$ and 7.5% below that of $(5, 7265)$; a seven-query algorithm needs $n \ge 602490$ to beat $(6, 90000)$. We also certify $(7, 650000)$ (0.362505), the first explicit seven-query algorithm.

- **Six queries search 91,000 elements**, which goes

TABLE I. Explicit exact ordered-search algorithms: k queries search n elements, which gives $(k/\log_2 n)\log_2 N$ queries for N elements. "Published" lists the largest n reported before this work, with the method used; "certified here" lists the largest n with an exact certificate in this work.

k	published n	$k/\log_2 n$	certified here	$k/\log_2 n$
3	56 (cited in [7])	0.516586	56	0.516586
4	605, SDP [7]	0.432863	605	0.432863
5	7265, LP [9]	0.389810	5000	0.406911
6	90000, SDP [10][a]	0.364572	**91000**	**0.364220**
7	–	–	**700000**	**0.360509**

[a] Numerical evidence; $N^*(6) < 94000$ is certified [10]. We also certify $(6, 90000)$ exactly.

beyond the numerical record. With the certified upper limit of Ref. [10], $91000 \le N^*(6) < 94000$.

- **Exact certificates of earlier numerical algorithms**: $(6, 90000)$ of Wu *et al.*, which turns their numerical evidence into a theorem, and $(4, 605)$, the optimum of CLP; also $(3, 56)$, $(5, 5000)$, $(6, 45000)$ and $(6, 80000)$.

The rest of the paper is organized as follows. Section II recalls the polynomial characterization and restates it in value form. Section III describes the certificate and the two checkers, Sec. IV how the chains were found, and Sec. V how a floating-point chain becomes an exact certificate. Section VI gives the results, Sec. VII the limits of the method and open problems, and Sec. VIII how the work was done.

II. THE POLYNOMIAL CHARACTERIZATION

A Laurent polynomial $q(z) = \sum_{|j|<n} c_j z^j$ is symmetric if $c_{-j} = c_j \in \mathbb{R}$; on the unit circle it is the cosine polynomial $q(\theta) = c_0 + 2\sum_{j=1}^{n-1} c_j \cos j\theta$. The Fejér kernel is $F_n(z) = \sum_{|j|<n}(1 - |j|/n)z^j$, so $F_n(\theta) = \sin^2(n\theta/2)/(n\sin^2(\theta/2))$.

Theorem 1 (FGGS [6]; CLP [7]; Theorem 8 of [9])**.** *An exact k-query translation-invariant quantum algorithm for n-element ordered search exists if and only if there are symmetric Laurent polynomials $q_0, \dots, q_k$ of degree less than n, nonnegative on $|z| = 1$, such that*

$$q_0 = F_n, \qquad q_k = 1, \qquad c_0(q_t) = 1 \quad (t = 0, \dots, k), \tag{1}$$

$$q_t(z) = q_{t-1}(z) \quad \text{if } z^n = (-1)^t \qquad (t = 1, \dots, k). \tag{2}$$

By Theorem 1 of Ref. [9], translation-invariant algorithms are optimal, so $N^*(k)$ is the largest n for which such a chain exists, and feasibility is monotone in n [9]. We do not construct the unitaries of the algorithms; the theorem guarantees that they exist.

Value form. Let $\theta_m = \pi m/n$, $m = 0, \dots, n$. The points $e^{\pm i\theta_m}$ with m even are the n roots of $z^n = 1$, and those with m odd are the n roots of $z^n = -1$. A symmetric Laurent polynomial of degree less than n is a cosine polynomial with $c_n = 0$; it is determined by its $n+1$ values $v_m = q(\theta_m)$, and by the type-1 discrete cosine transform

$$c_0 = \frac{1}{n}\sum_{m=0}^{n} w_m v_m, \qquad c_n = \frac{1}{n}\sum_{m=0}^{n} w_m (-1)^m v_m = 0, \tag{3}$$

with $w_0 = w_n = \frac{1}{2}$ and $w_m = 1$ otherwise. So $c_0 = 1$ and $c_n = 0$ hold exactly when the values on each parity class average to one over the corresponding n roots of unity. Conditions (2) say that q_t and q_{t-1} share their values on one parity class, alternating with t. A chain is therefore a path of "half-vectors"

$$u_0 \to u_1 \to \cdots \to u_{k-1}, \tag{4}$$

where u_t lives on the parity class of $t+1$ (odd m for u_0), u_0 is the odd-m part of F_n, $u_0(m) = 1/(n\sin^2(\pi m/2n))$, $u_{k-1} = 1$, every u_t has mean one, and q_t interleaves u_{t-1} and u_t (q_0 itself is n at $m = 0$ and 0 at the other even m). Conditions (1) and (2) then hold by construction, and the problem is to choose $u_1, \dots, u_{k-2}$ so that every interleaved q_t is nonnegative on the whole circle, not only at the samples. All of the difficulty lies between the samples: the interleaving makes q_t oscillate at frequency close to n, with narrow dips between neighboring samples.

III. CERTIFICATES AND CHECKERS

A certificate lists, for $t = 1, \dots, k-1$, the integer coefficients of q_t over a common denominator D divisible by n, and the integer coefficients of a polynomial $h_t(z) = \sum_{j=0}^{n-1} h_{t,j} z^j$ over a denominator H. The verifier fixes $q_0 = F_n$ and $q_k = 1$ and checks in exact integer arithmetic:

1. every q_t has constant term 1;
2. *root conditions.* On $z^n = s$ we have $z^{-j} = s\, z^{n-j}$, and $1, z, \dots, z^{n-1}$ are linearly independent as functions on those n roots. So with $d = q_t - q_{t-1}$, condition (2) holds if and only if $d_j + s\, d_{n-j} = 0$ for $j = 1, \dots, n-1$, where $s = (-1)^t$;
3. *nonnegativity.* The residual $r_t = q_t - |h_t|^2$, a symmetric Laurent polynomial, must be diagonally dominant: $r_{t,0} \ge 2\sum_{j\ge1} |r_{t,j}|$.

The last condition implies $q_t(\theta) \ge |h_t(e^{i\theta})|^2 + r_{t,0} - 2\sum_{j\ge1}|r_{t,j}| \ge 0$ for all θ. We call $(r_{t,0} - 2\sum_{j\ge1}|r_{t,j}|)$ the certified margin of q_t. A positive margin is always achievable: by the Fejér–Riesz theorem [12, 13] a strictly positive $q_t - \gamma$ equals $|h|^2$ for a polynomial h of degree $n-1$, and rounding h perturbs $|h|^2$ by much less

than γ. The autocorrelation $|h_t|^2$ is computed exactly by Kronecker substitution, that is, by one multiplication of two large integers. The verifier (`verifier/verify.py`, Python standard library only) is pinned by a hash, version 1.1. Besides the verdict it raises an alarm if a certificate contradicts a proven or published bound: $k < \ln(n)/\pi - 2$ (the bound of Ref. [4], weakened), or n above the published feasibility limits 6, 605, 7265 and 93999 for $k = 2, 4, 5, 6$. Either alarm would point to a bug in the verifier or an error in the literature. Its test suite contains the certificates for $(3, 56)$ and $(4, 605)$, every verdict path, 17 malformed or tampered inputs, and cross-checks of the Kronecker autocorrelation and of the root conditions against brute force.

An independent checker shares no code with the verifier and uses a different route for each condition, in rational arithmetic from FLINT [14] via python-flint. For the root conditions it checks that $z^n - s$ divides the integer polynomial $z^{n-1}(q_t - q_{t-1})(z)$; it forms $|h_t|^2$ as the product $h_t(z) \cdot z^{n-1}h_t(1/z)$ and repeats the diagonal-dominance test. On all nine certificates the two checkers agree to the last printed digit on every margin. An audit script records both results with the verifier's hash, and adds a floating-point evaluation of every q_t on a grid 16 times finer than n as a sanity check.

IV. FINDING THE CHAINS

Reduced basis. Each free half-vector u_t ($t = 1, \ldots, k-2$) is parametrized by its individual samples for $m \le m_0 = 100$ and, beyond, by piecewise-linear hat functions in $\log\theta$ with knots spaced by a ratio of at most $\rho = 1.05$ up to $\theta = \pi$. This gives 708 variables at $(6, 45000)$, 766 at $(6, 90000)$ and $(6, 91000)$, and 1158 at $(7, 650000)$ and 1168 at $(7, 700000)$.

Weighted max–min. The natural scale of the chain is the Fejér tail, so we maximize the weighted margin $\min_{t,\theta} q_t(\theta)/\omega(\theta)$ with $\omega(\theta) = \min\{1, 1/(n\sin^2(\theta/2))\}$. The minimum is replaced by the smooth $\mathrm{softmin}_\mu(x) = -\mu \log \sum_i e^{-x_i/\mu}$, which lies within $\mu \ln(\#\text{points})$ of the true minimum, and is taken over two sets of points: a grid 8 times finer than n (evaluated by batched type-1 discrete cosine transforms), and "extra points" at the local minima of q_t found in earlier passes, which catch the narrow dips that the grid misses. Samples where u_t is fixed (those of u_0 and $u_{k-1} = 1$) are masked: on the samples of u_0 the ratio is exactly $q/\omega = 1$, which would cap the objective. We maximize with L-BFGS-B [15] and continuation in μ from 1 to 0.003.

Passes. After each optimization, every local minimum of every q_t on a 16-fold grid is refined by five Newton steps, with q_t' and q_t'' interpolated from their own oversampled grids. The minima that fall below the current objective become extra points, and the pass is repeated until none is added. The smallest refined value of q_t/ω over all t is the "true minimum" of the chain; a chain is handed to the certificate builder when it is positive. An earlier version refined only the minima whose parabolic estimate from the grid looked low. At large θ, where $\omega \sim 1/n$ is smaller than the error of that estimate, it skipped a band of 28 real negative minima in the $(6, 90000)$ chain. The certificate builder detected them, because the chain was negative on its finer grid, and after the fix every grid minimum is refined.

Large n. For $k = 7$ three changes keep a pass near $O(n \log n)$: extra points are sparse rows that interpolate the oversampled grid with a Gaussian-windowed sinc, in the spirit of nonuniform FFTs [16], instead of dense rows of length n; local minima are refined from interpolated derivative grids; and the $k-1$ transforms of one evaluation run as one threaded batch, 0.6 s per objective and gradient at $n = 650000$. The decisive change was a log parametrization: every block is written as $a = n\,e^x/(r \cdot e^x)$, where $r \cdot a = n$ is the mean-one constraint. It keeps $a > 0$, makes the constraint exact and turns optimizer steps into relative changes, which matters because the samples of a chain span many decades. With a linear diagonal scaling instead, the weighted grid minimum after 1000 iterations at $\mu = 1$ was -800; with the log parametrization it was -53.

Warm starts. At fixed sample index m (not fixed θ) the shape of a chain changes slowly with n. When the basis has the same number of variables at two sizes, we copy the coefficients, rescale each block to mean one, and map extra points by $\theta \mapsto \theta\, n_{\mathrm{old}}/n_{\mathrm{new}}$. The $(6, 90000)$ chain moved to $n = 91000$ this way starts with true minimum -3.4, against -531 after a cold first pass. When the number of variables changes, as from $(7, 650000)$ (1158) to $(7, 700000)$ (1168), each coefficient beyond m_0 is stored as a ratio to the envelope $U^{1-t/(k-1)}$, where $U(\theta) \approx 1/(n\sin^2(\theta/2))$ is the envelope of u_0, and carried over at the same fraction of the knot range (fixed m at the bottom of the basis, fixed θ at $\theta = \pi$). The $(7, 700000)$ start then has true minimum -129, against -874 when the samples are copied at fixed m.

Table II lists the runs behind the certificates; all ran on a laptop (Apple M4, 16 GB of memory).

At $(7, 650000)$ the weighted minimum over the grid and the extra points crossed zero in pass 6. The true minimum became positive in pass 17 and reached 0.987 in pass 21, by which time 1.65 million exact local minima served as constraints. Warm-started from that chain, the run at $(7, 700000)$ had a positive minimum over the grid and the extra points already at the first continuation stage $\mu = 0.1$ (cold at 650000: -5.0), and reached a true minimum of 0.990 after 6 passes and 643,058 extra points.

V. FROM A FLOATING-POINT CHAIN TO AN EXACT CERTIFICATE

Exact rounding. The coefficients are rounded to the denominator $D = n \cdot 2^{57}$ step by step. Each increment $q_t - q_{t-1}$ is projected in integers onto the subspace allowed by (2), so the root conditions hold exactly. The last

TABLE II. Search runs. "Passes" counts optimization passes, and "true min" is the smallest refined local minimum of q_t/ω over all t, recomputed for every certified chain with the corrected minimum check of Sec. IV. The $(6, 45000)$ and $(6, 80000)$ chains were found by an earlier dense implementation of the same search. Two unsuccessful runs are listed at the bottom.

k	n	variables	passes	time	true min
6	45000	708	6	13 min	0.994
6	80000	754	12	2.2 h	0.808
6	90000	766	40 + 40	1.8 h + 84 min[a]	0.265
6	91000	766	14	35 min[b]	0.147
7	650000	1158	21	16.1 h	0.987
7	700000	1168	6	3.8 h[b]	0.990
6	92000	766	40	1.9 h	−0.17
7	850000	1188	4	7 h	≈ -25[c]

[a] The run was resumed after the minimum check was fixed (Sec. IV).
[b] Warm start from the chain in the row above.
[c] Weighted grid minimum, the same on every q_t; a basis twice as fine (2200 variables) reached −23.8 in two further passes.

polynomial q_{k-1} is not rounded but solved for: its values on both parity classes are fixed, by q_{k-2} and by $q_k = 1$.

Spectral factors. For each q_t we take γ_t to be half the minimum of q_t on a fine grid and compute the minimum-phase spectral factor h_t of $q_t - \gamma_t$, rounded to the denominator $H = 2^{60}$. We start from the cepstrum [17] on a half-spectrum FFT grid of M points, apply truncated Newton steps of Wilson's method [18], and double M up to 2^{26} until $1/|h_t|^2$ is resolved. If the factor error is still above the budget left by γ_t, mixed-precision Newton refinement follows, with the residual $q_t - |h_t|^2$ computed exactly by FLINT.

Spikes. At $n = 650000$ the roots of $q_t - \gamma_t$ come within 1.4×10^{-8} of the unit circle, closer than the grid spacing $2\pi/2^{26} = 9.4 \times 10^{-8}$. Then $1/|h_t|^2$ has spikes narrower than the grid, and plain Newton refinement oscillates: the first attempt failed with a certified margin of -7.7×10^{-7} on q_3, and both checkers rejected it. The fix computes the Fourier coefficients of the Newton step exactly near each spike. Around every minimum where the estimated root distance $\sqrt{2(q_t - \gamma_t)/q_t''}$ is below 3.5 grid spacings, a smooth partition of unity isolates the spike. The spike's integral is computed by a tan-mapped Gauss rule on the core and Gauss panels on the transition, minus the grid's own share. The corrections are then folded into all n coefficients by a 26-term moment expansion and the transpose of the windowed-sinc interpolation. On a test chain this cut the coefficient error from 0.22 to 3×10^{-10}. With it, q_5 (18,304 spikes) converged in two steps and q_6 (56,639 spikes, plain-FFT error 3.7×10^{-6}) in eight. The certificate took 22 minutes and 6.6 GB to build. The $(6, 91000)$ certificate needed the same correction for 32,504 spikes in q_5, with roots 1.2×10^{-8} from the circle, and took 3 minutes.

TABLE III. The nine certificates. The margin is the smallest certified margin over t (absolute, in units of the coefficients of q_t); both checkers report the same value. Times are for the standard-library verifier and the FLINT checker on the laptop of Table II (for $(7, 700000)$ the FLINT checker ran alongside another search); files are gzipped JSON except the two smallest.

k	n	$k/\log_2 n$	margin	size	verifier	FLINT
3	56	0.516586	2.4×10^{-3}	4.8 kB	< 0.1 s	< 0.1 s
4	605	0.432863	4.3×10^{-5}	74 kB	< 0.1 s	< 0.1 s
5	5000	0.406911	1.0×10^{-4}	0.35 MB	0.2 s	< 0.1 s
6	45000	0.388158	1.1×10^{-5}	3.9 MB	5.5 s	0.5 s
6	80000	0.368376	6.1×10^{-6}	6.8 MB	13 s	1.0 s
6	90000	0.364572	1.7×10^{-6}	7.6 MB	18 s	1.4 s
6	91000	0.364220	9.3×10^{-7}	7.7 MB	21 s	1.6 s
7	650000	0.362505	7.5×10^{-7}	64.4 MB	7.1 min	12 s
7	700000	0.360509	7.1×10^{-7}	68.8 MB	8.0 min	19 s

A finer starting grid. At $(7, 700000)$ the correction was not enough for q_6: its 70,223 spikes come within 8.7×10^{-9} of the circle, the cepstrum on 2^{26} points left a factor error of 1.2×10^{-5}, and the first Newton step overshot at the spikes. From there the residual fell only by a factor of 4 per step, as for Newton's method on $x^2 = a$ started at $x \gg \sqrt{a}$, and after 10 steps it was 1.7×10^{-6}. Both checkers rejected the certificate (margin −0.33 on q_6). Recomputing that one factor on 2^{27} points gave a starting error of 2.6×10^{-11}, and two refinement steps reached 7.5×10^{-16} (6 minutes, with the other factors kept). The builder now retries on a grid twice as fine whenever the refinement does not converge.

VI. RESULTS

Table III lists the certificates and Fig. 1 compares the constants with the published ones and with the asymptotic optimum.

Seven queries. The certificate for $(7, 700000)$ proves that ordered search on N elements can be solved exactly with $7\lceil \log_{700000} N \rceil$ queries, $0.360509 \log_2 N$ asymptotically. To our knowledge this is the best explicit exact ordered-search algorithm; its constant is still a factor 1.63 above the optimum $\ln 2/\pi$ [4, 11]. The certificate for $(7, 650000)$, found first, gives $0.362505 \log_2 N$.

Six queries. The certificate for $(6, 91000)$ goes beyond the numerical record 90000 of Ref. [10]. Together with their certified infeasibility at 94000 and monotonicity, $91000 \le N^*(6) < 94000$. The certificate for $(6, 90000)$ makes their numerical result rigorous. The sequence of six-query certificates 45000, 80000, 90000, 91000 shows the margin shrinking as n approaches the limit of the search (Table III).

Smaller k. The $(4, 605)$ certificate confirms exactly that four queries suffice for 605 elements, the optimum

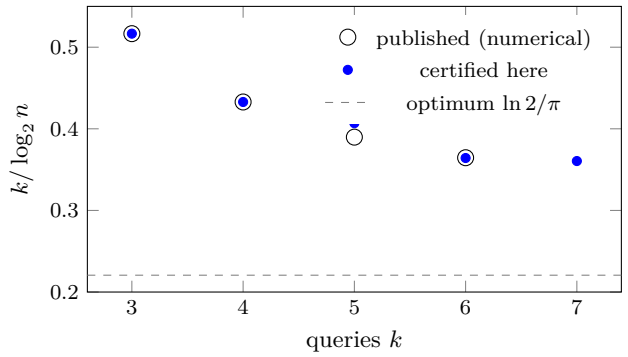


FIG. 1. Constants $k/\log_2 n$ of explicit exact algorithms: the largest published n for each k [7, 9, 10] and the largest certified here. The dashed line is the asymptotic optimum $\ln 2/\pi = 0.2206$ [4, 11].

of CLP [7]. The cutting-plane linear program we used for calibration reproduces the known thresholds numerically: two queries 6 but not 7, three queries 56 but not 57, four queries 605 but not 606 (optimal weighted margin $+8.7\times 10^{-5}$ at 605 and -4.2×10^{-4} at 606). The reduced basis is nearly lossless at these sizes: it is feasible at exactly 605 and infeasible at 606, and at $k = 5$ its best relative margin at the true limit 7265 is -0.004. We did not certify $(5, 7265)$; the $(5, 5000)$ certificate is a calibration.

VII. LIMITS AND OPEN PROBLEMS

How far the search reaches. At $k = 6$ the same search ended at a true minimum of -0.17 for $n = 92000$ after 40 passes, with the minimum over grid and extra points still rising slowly. Its limit is therefore near 91500–92000, within about 3% of the upper limit 94000. The true minima of the certified six-query chains show how the margin behaves near the limit (Table II). Far from the limit it stays close to its ceiling, $q_1/\omega = 1$ at the fixed samples of u_0 (0.994 at $n = 45000$); nearer the limit it falls steeply (0.808 at 80000, 0.265 at 90000, 0.147 at 91000). The seven-query chains still have 0.987 at 650000 and 0.990 at 700000, so even 700000 is likely far from $N^*(7)$. At $k = 7$ the gap is larger. The ratios $N^*(k)/N^*(k-1)$ are 9.3, 10.8, 12.0 and about 12.5 for $k = 3, \ldots, 6$, and must eventually approach $e^\pi \approx 23.1$, which suggests $N^*(7) \approx 1.1$–1.2 million. A run at $n = 850000$ stalled at a weighted grid minimum of about -25, the same on every q_t and at every frequency. Refitting that chain in a basis twice as fine ($m_0 = 200$, $\rho = 1.025$, 2200 variables) moved the minimum only to -23.8 in two passes. So the resolution of the basis is not what limits the search; the optimizer is. In the coefficients the restricted problem is a concave max–min, since each q_t is linear in them, and the softmin is within $\mu \ln(\#\text{points}) \approx 17.5\mu$ of its optimum. Yet in the cold first pass at $n = 850000$ every stage of the continuation ended far below that bound (at $\mu = 1$, -106, although -22 is attainable). A solver that converges faster on the restricted linear program, for instance stabilized cutting planes with an interior-point method, is one route toward $n \geq 850000$. The other is a ladder of warm starts, which took the present solver from 650000 to 700000 in 3.8 hours with the margin still at its ceiling.

Size of the certificates. A certificate stores $2(k-1)n$ large integers, so the $(7, 700000)$ file has 69 MB compressed and takes 8 minutes to check with the standard library. Larger k will need either a more compact format, for instance coefficients generated from a short description, or a certificate that exploits the structure of the chain.

Open problems. (i) Close the gap $91000 \leq N^*(6) < 94000$, and certify the five-query optimum $(5, 7265)$, which sits at the limit of feasibility, so its margins are tiny. (ii) Certify infeasibility exactly, for example at $(5, 7266)$ or $(6, 94000)$, with a rational dual certificate. (iii) Explicit finite instances of the Carolan–Childs construction [11], to compare its finite-n performance with $N^*(k)$. (iv) Explicit exact algorithms that match the zero-error algorithm of Ben-Or and Hassidim [8], whose $0.32 \log_2 N$ expected queries are still below every exact constant here.

VIII. HOW THE RESULTS WERE FOUND

The problem was set up on 28–29 October 2026 by a language-model session (Claude Opus 5.5, Anthropic), after a human operator asked it to start a new problem from a list of verifiable open problems and solve it. The session wrote the verifier, its tests and the record snapshot, and then carried out the search alone. The literature check that found Refs. [10, 11] came after the pipeline was built. The verifier was then revised once, to version 1.1: the record it compares against changed from $(5, 7265)$ to the numerical $(6, 90000)$, and the alarm limit 93999 for $k = 6$ was added. All certificates were audited under version 1.1 with matching hashes. The independent checker was written by the same session, but with a different route for each condition and a different arithmetic library. The certificates for $k \leq 6$ up to $n = 80000$ were audited on 29 October, $(7, 650000)$ on 30 October after a 16-hour search, and $(6, 90000)$, $(6, 91000)$ and $(7, 700000)$ on 31 October. The human operator chose the problem, asked for the seven-query run and for larger n before a paper was drafted, and decided when to commit results. This manuscript was drafted by the same model from the results summary, the audit records and the run logs, and checked against them.

ACKNOWLEDGMENTS

We thank Y. Wu, H. Liu, W. Gao, Y. Su, T. Li, D. Ge and Y. Ye for the numerical frontier and the infeasibil-

ity certificate of Ref. [10], and J. Carolan, A. M. Childs, M. Kovacs-Deak and L. Schaeffer for the characterization of Ref. [9] on which the verifier rests. The certificates were found and checked by Claude Opus 5.5 (Anthropic), which also wrote the verification software and this manuscript.

DATA AVAILABILITY

The nine certificates, the frozen verifier with its tests and hash pin, the independent checker, the audit records and the SHA-256 sums of all certificates will be released with this paper at `https://github.com/maxxfan/aixmathxphysics/releases/tag/ordered-search-v1`. The seven-query certificates (64.4 and 68.8 MB) are distributed as release assets. Running `python3 verifier/verify.py` on a certificate re-checks it with the Python standard library, and `results/check_independent.py` (python-flint) rechecks it independently. The search and certificate-building programs are available on request.

[1] H. Buhrman and R. de Wolf, A lower bound for quantum search of an ordered list, Inf. Process. Lett. **70**, 205 (1999).

[2] E. Farhi, J. Goldstone, S. Gutmann, and M. Sipser, A limit on the speed of quantum computation for insertion into an ordered list, arXiv:quant-ph/9812057 (1998).

[3] A. Ambainis, A better lower bound for quantum algorithms searching an ordered list, in *Proceedings of the 40th IEEE Symposium on Foundations of Computer Science* (1999), p. 352, arXiv:quant-ph/9902053.

[4] P. Høyer, J. Neerbek, and Y. Shi, Quantum complexities of ordered searching, sorting, and element distinctness, Algorithmica **34**, 429 (2002), arXiv:quant-ph/0102078.

[5] A. M. Childs and T. Lee, Optimal quantum adversary lower bounds for ordered search, in *Automata, Languages and Programming (ICALP 2008)*, Lecture Notes in Computer Science **5125**, 869 (2008), arXiv:0708.3396.

[6] E. Farhi, J. Goldstone, S. Gutmann, and M. Sipser, Invariant quantum algorithms for insertion into an ordered list, arXiv:quant-ph/9901059 (1999).

[7] A. M. Childs, A. J. Landahl, and P. A. Parrilo, Quantum algorithms for the ordered search problem via semidefinite programming, Phys. Rev. A **75**, 032335 (2007), arXiv:quant-ph/0608161.

[8] M. Ben-Or and A. Hassidim, Quantum search in an ordered list via adaptive learning, arXiv:quant-ph/0703231 (2007).

[9] J. Carolan, A. M. Childs, M. Kovacs-Deak, and L. Schaeffer, Translation-invariant quantum algorithms for ordered search are optimal, ACM Trans. Quantum Comput. **7**, no. 3 (2026), arXiv:2503.21090.

[10] Y. Wu, H. Liu, W. Gao, Y. Su, T. Li, D. Ge, and Y. Ye, Matrix-free GPU semidefinite programming for quantum ordered search at the $k = 6$ frontier, ICML 2026 (poster), `https://openreview.net/forum?id=GSGOyxRrBo`.

[11] J. Carolan and A. M. Childs, Optimal quantum algorithms for ordered search, arXiv:2609.34293 (2026).

[12] L. Fejér, Über trigonometrische Polynome, J. Reine Angew. Math. **146**, 53 (1916).

[13] F. Riesz, Über ein Problem des Herrn Carathéodory, J. Reine Angew. Math. **146**, 83 (1916).

[14] The FLINT team, FLINT: Fast Library for Number Theory, `https://flintlib.org`; python-flint, `https://github.com/flintlib/python-flint`.

[15] D. C. Liu and J. Nocedal, On the limited memory BFGS method for large scale optimization, Math. Program. **45**, 503 (1989).

[16] A. Dutt and V. Rokhlin, Fast Fourier transforms for nonequispaced data, SIAM J. Sci. Comput. **14**, 1368 (1993).

[17] A. H. Sayed and T. Kailath, A survey of spectral factorization methods, Numer. Linear Algebra Appl. **8**, 467 (2001).

[18] G. T. Wilson, Factorization of the covariance generating function of a pure moving average process, SIAM J. Numer. Anal. **6**, 1 (1969).